

# PostgreSQL Database Extensions Optimization

Alex Osborne

Undergraduate Research

Professor Grant Scott

## **Abstract**

This project explores the optimization of PostgreSQL database extensions, originally developed in C, by transitioning them to Rust to enhance security and performance. The choice of Rust was driven by its safety features and efficiency, as well as its compatibility with PostgreSQL through the pgrx framework. The endeavor involved learning Rust, analyzing existing C based extensions, and beginning the translation of a key CPU function optimization module. This translation aimed to make use of Rust's performance advantages and robust type safety, crucial for the reliability and speed of database operations. The complete translation was not achieved, however, substantial progress was made in understanding and partially implementing the extensions in Rust. This experience highlighted the complexities of adapting existing software to new technologies and underscored the potential of Rust in database extension development.

## **Introduction**

PostgreSQL is a powerful, open source object-relational database system known for its community support, scalability, and ability to handle large volumes of data. It is used in various industries and scenarios for complex data management needs. One of the strengths of PostgreSQL is its extensibility, which allows developers to introduce custom functions and

optimizations through extensions. These extensions can enhance the database's functionality and performance but are often limited by the language in which they are written. In this project, the focus is on extensions originally developed in C, which are utilized for optimizing specific database operations.

While C has been a staple in systems programming due to its speed and close to metal nature, it poses several challenges, particularly in terms of security and maintainability. Common issues include memory management errors and buffer overflows, which can compromise the security and stability of database systems. Additionally, C's manual memory management can lead to errors that are hard to detect and resolve. These challenges necessitate exploring other programming languages that maintain the efficiency of C but improve safety and developer productivity.

The primary objective of this research is to explore the feasibility and benefits of transitioning existing PostgreSQL extensions from C to Rust. Rust offers memory safety guarantees through its ownership model, alongside performance on par with C, making it an appealing choice for system level programming in a database context. This project aims to:

- Analyze the current C based extensions for CPU function optimization to identify potential areas for improvement.

- Begin the process of translating these extensions to Rust to enhance their security and performance.

## **Research**

Many projects have used languages like C to harness hardware accelerators within PostgreSQL. Notably, a previous graduate student's project demonstrated how CUDA could be

integrated to optimize various functions like vector addition and heart rate estimation through parallel reduction algorithms executed on GPUs. These extensions highlight PostgreSQL's flexible architecture that allows for significant performance improvements via external hardware acceleration.

Over the years, PostgreSQL has seen various enhancements in its extensibility. These developments have been primarily driven by the need to optimize database operations, manage large datasets more efficiently, and secure data handling. Extensions, often written in C due to its performance advantages, have played a crucial role in this evolutionary process, providing custom solutions to specific database challenges.

Introduced by Mozilla, Rust has gained attention for its ability to provide memory safety without sacrificing performance. Studies have shown Rust's effectiveness in achieving safety through its ownership and borrowing system, which statically prevents bugs common in systems programming, such as data races and buffer overflows (Bugden and Alahmar, 2022). This makes Rust a promising candidate for developing robust and secure database extensions.

Extensive research has compared Rust with traditional systems languages like C and C++. Findings suggest that Rust not only matches but in some cases surpasses these languages in performance while significantly elevating safety standards. For instance, benchmarking tests reveal Rust's competitive edge in managing memory usage and CPU time efficiently, critical metrics for database operations (Bugden and Alahmar, 2022).

The design principles of Rust, focusing on zero cost abstractions and memory safety, make it particularly suitable for environments where concurrency and security are crucial. This aligns with the needs of modern databases that handle concurrent transactions and secure data processing without compromising performance.

In conclusion, the transition to Rust for developing PostgreSQL extensions is supported by its proven capabilities in similar programming environments. The groundwork laid by previous studies and projects provides a solid basis for exploring Rust's application in optimizing database extensions.

## **Methodology**

This study attempted a systematic approach at exploring the possibilities of transitioning PostgreSQL extensions from C to Rust. The primary focus was to analyze existing C based extensions, identify their inherent inefficiencies and security vulnerabilities, and subsequently translate them to Rust.

The methodology was based on comparative analysis, where specific functions within the extensions were isolated, reviewed, and reengineered in Rust. The criteria for analysis were determined based on performance metrics, memory usage, and security loopholes that are typical in C based systems programming.

The translation process was designed to be iterative, allowing for the gradual refactoring of code. This would create an opportunity for each function to be benchmarked and compared to its previous version. Rust's functionality, particularly the memory safety features, were implemented to reconstruct the architecture of these functions.

## **Tools and Technology**

Software and Programming Languages:

PostgreSQL: The latest stable release of PostgreSQL was used as the primary database management system framework for the extensions.

Rust: Employed for rewriting the extensions, again utilizing the most recent version which includes the latest safety and performance features.

C: The original programming language of the database extensions, served as the starting baseline of the project.

#### Development Tools:

Cargo: Rust's package manager and compiler, facilitated dependency management and streamlined the build process for Rust components.

pgrx: A framework that aids in creating PostgreSQL extensions in Rust, used to integrate the newly written Rust code seamlessly.

#### Hardware:

A personal, high performance, windows laptop computer as it was the most convenient and readily accessible machine. Equipped with an 11th Gen Intel(R) i7-11800H processor, 16 GB of RAM, 1 TB Samsung SSD 970 EVO Plus, and an NVIDIA GeForce RTX 3060 Laptop GPU. This setup is capable of simulating testing locally and reflects the performance characteristics of similar research study hardware.

This methodology, with its focus on critical analysis, highlights the research's objective to not only translate but enhance the extensions. Thus ensuring that the transition to Rust translates to real improvements in performance and security.

## **Implementation**

As previously outlined, the initial stage focused on assessing the feasibility and strategic planning for translating critical CPU function optimization modules from C++ to Rust, aimed at

enhancing security and performance. Due to time constraints and other obligations, the plan was to pace the development in a way that would allow it to be picked up and put down frequently.

The core algorithms, such as energy computation, smoothing functions, peak detection, and filtering, were methodically translated over time. Special attention was given to utilizing Rust's strong type system and memory safety features without sacrificing efficiency.

**Compute Energy Function:** The `compute_energy_cpu` function was redesigned to make use of Rust's ability to handle generics with trait bounds, ensuring that the function remained flexible and type safe.

**Smoothing Algorithm:** Translating the smoothing function involved rethinking the convolution operation to make effective use of Rust's iterators and functional programming capabilities.

**Peak Detection and Localization:** The translation of these functions required mindfulness of indexing and bounds checking, areas where Rust's safety features really shine.

**Butter Lowpass Filter:** The digital filter implementation was adapted to Rust's zero cost abstraction model. The filter function was reimplemented using Rust's arrays and slices, ensuring memory safety without the overhead.

## Challenges

Rust's learning curve was steep, particularly in understanding ownership, borrowing, and lifetimes. When deciding on the translation language, prior experience was taken into consideration but did not outweigh the potential benefits Rust could provide. Starting from scratch proved to be very difficult, especially considering the complex nature of a multi paradigm language. This was initially a significant hurdle in translating the source code and specifically pointer arithmetic directly into safe Rust code.

While Rust's safety features are advantageous, they require different optimization strategies compared to procedural languages. Performance tuning was necessary to match the expected speed of the C++ code, especially for the heavily used and computationally intense functions.

Translating C++ templates into Rust generics involved a deep understanding of both language type systems. In the initial phases of redevelopment it was easier to completely ignore the generic inputs and outputs and instead focus on the meat of the function. Rust's trait bounds provided a powerful, albeit very initially confusing, way to handle generic programming.

Towards the beginning of the semester, naive optimism of the project's timeline was quickly corrected. It seemed the process of changing a few C++ files to a new programming language would be easy but it turned out to be much more time consuming. The difficulty of learning a new programming language on its own was severely underestimated. Balancing coursework and other academic responsibilities in a challenging semester made time constraints all the more apparent.

The endeavor to translate CPU function optimization modules from C++ to Rust has been both challenging and enlightening. This project not only required a deep dive into the nuances of Rust's programming paradigms but also tested resilience and adaptability in managing time and expectations. Through an iterative process, each function was reconceptualized to leverage Rust's robust safety features and efficient memory management. The experience beautifully showcased the importance of realistic project planning as well as accounting for unexpected difficulties. Moving forward, the insights gained from this project will undoubtedly influence future coding practices and architectural decisions.

## Reflection

The journey through this project was not just about translating PostgreSQL extensions from C to Rust. It served as a way to experience the practical application of concepts learned by analyzing accredited research. This has significantly broadened my understanding of programming languages and also language translation efforts. Since more companies continue to explore Rust as an option for their codebase, I can't help but feel lucky for the path this project took.

Looking back to the beginning of the semester, the first notable lesson was the importance of choosing the right tool for the right job. Rust's safety features and modern paradigms provided a capable framework that, while challenging, brought a new level of reliability and efficiency to the project. Learning the language from scratch made me appreciate its memory safety abilities, especially in an environment where data integrity is crucial.

The project also taught me the value of resilience and adaptability in research, something I had not experienced before. There were a number of challenges, from steep learning curves in Rust, balancing academic loads, to even losing progress due to technical faults. It was a reminder that research is often unpredictable and demands a proactive approach to problem solving and learning.

Furthermore, this undertaking highlighted the importance of effective communication and collaboration, even in primarily solo projects. Seeking help from peers, discussing problems with mentors, and engaging with developer communities are all crucial to overcome obstacles. Had this aspect of the project been more strongly utilized, potentially much more progress could have been made. Such a pitfall serves as a reminder to build and nurture a strong and supportive collaborative network.



Professionally, this project has prepared me to tackle more complex software development challenges in the future. The hands on experience with Rust and reformulation of existing code has undoubtedly primed me for real world scenarios.

While the project's technical goals were only partially met, the educational and professional growth I experienced exceeded expectations. The lessons learned from this project will influence my future projects and career path, steering me towards more secure, efficient, and reliable programming practices.

## **Future Work**

The completion of this initial phase has laid out the foundation for continued development and exploration. The next steps involve not only completing the translation of remaining database modules but also refining those that have already been ported. The progress achieved so far provides a strong basis for future work, where more complex aspects of the database can be reenvisioned within Rust's paradigms.

One major focus would be on enhancing the performance of the already translated modules. Now that a basic translation has been established, the next phase will include performance tuning and optimization to ensure that these modules not only match but ideally exceed the performance of the original C++ implementations. This would involve an even deeper dive into Rust's advanced features such as concurrent programming.

Upcoming semesters, while still intimidating, are anticipated to not have such demanding workloads. With the foundational knowledge of Rust now firmly in place, efforts can be directed towards more ambitious challenges and innovations. The learning curve will no longer be as

steep, enabling a quicker development pace and more time for refining and expanding the project's scope.

The experience gained has also highlighted the importance of collaboration in complex projects. Moving forward, I plan to actively seek out and engage more with peers, faculty, and the Rust community to help overcome technical challenges and find solutions. With a working prototype already in hand, it will be easier to demonstrate the tangible progress of the project, making collaboration more appealing and effective.

Lastly, documenting the development process and findings will be prioritized to aid in future efforts by myself or others. Should there be a meaningful success story, publications and presentations at relevant forums could be on the agenda. This would share the knowledge gained and invite feedback to further refine the approach and outcomes.

The future stages of this project are not just about continuing but expanding the scope and effect of Rust in database systems. With a clear roadmap and the hard earned experience of the past semester, the project is in a position to make strides in the optimization and functionality of PostgreSQL extensions.

## References

Alahmar, A., & Budgen, W. (2022). Rust: The Programming Language for Safety and Performance. <https://doi.org/10.48550/arXiv.2206.05503>

Tripuramallu, D., & Singh, S., & Deshmukh, S., & Pinisetty, S., & Shivaji, S., & Balusamy, R., & Bandeppa, A. (2024). Towards a Transpiler for C/C++ to Safer Rust. <https://doi.org/10.48550/arXiv.2401.08264>